

# CS 4400

## Computer Systems

---

### LECTURE 18

Virtual memory

Main memory as a cache

Address translation

# Virtual Memory

---

- ***Virtual memory*** (VM) is an abstraction of main memory
- Treats main memory (MM) as a cache for disk
  - Transfers data back and forth between disk and MM, as needed
- Simplifies memory management
  - provides each process with a uniform address space

# Virtual Memory

---

- Protects the address space of each process
  - prevents some process from inadvertently writing to the memory used by another process
- Works silently and automatically
  - *without any intervention* from the application programmer

# Why Care About VM

---

- VM pervades all levels of computer systems
  - hardware exceptions, linkers, loaders, processes, ...
- VM gives applications powerful capabilities.
  - create/destroy chunks of memory, map chunks of memory to portions of disk, share memory with other processes

# Why Care About VM

---

- VM, used improperly, can lead to difficult bugs
  - any variable reference, pointer dereference, or `malloc` call uses VM
  - possible behaviors: “Segmentation fault”; long, silent run before crashing; run to completion with incorrect results

# Physical Addressing

---

- MM is an array of  $M$  contiguous byte-sized cells
  - each byte has a unique ***physical address*** (PA)
- ***Physical addressing*** is using the PA to access memory
  - used by early PCs, embedded microcontrollers, and others

# Physical Addressing

---

- Executes a load instruction
  - Generate a PA  $x$
  - Passes  $x$  to main memory (via memory bus).
- Main memory
  - Fetches the word starting at PA  $x$
  - Returns it to the CPU
  - Stores in a register

# Virtual Addressing

---

- Another form of addressing uses a ***virtual address*** (VA).
- ***Virtual addressing*** is using the VA to access memory.
  - used by modern processors for general-purpose computing



# Virtual Addressing

---

- On load:
  - CPU generates a VA  $y$
  - $y$  is converted to the appropriate PA  $x$
  - Passed to main memory
- The translation of VA to PA requires close cooperation between the CPU and OS
  - memory management unit (MMU)—dedicated hardware for translating VAs, using look-up table stored in memory

# Address Space

---

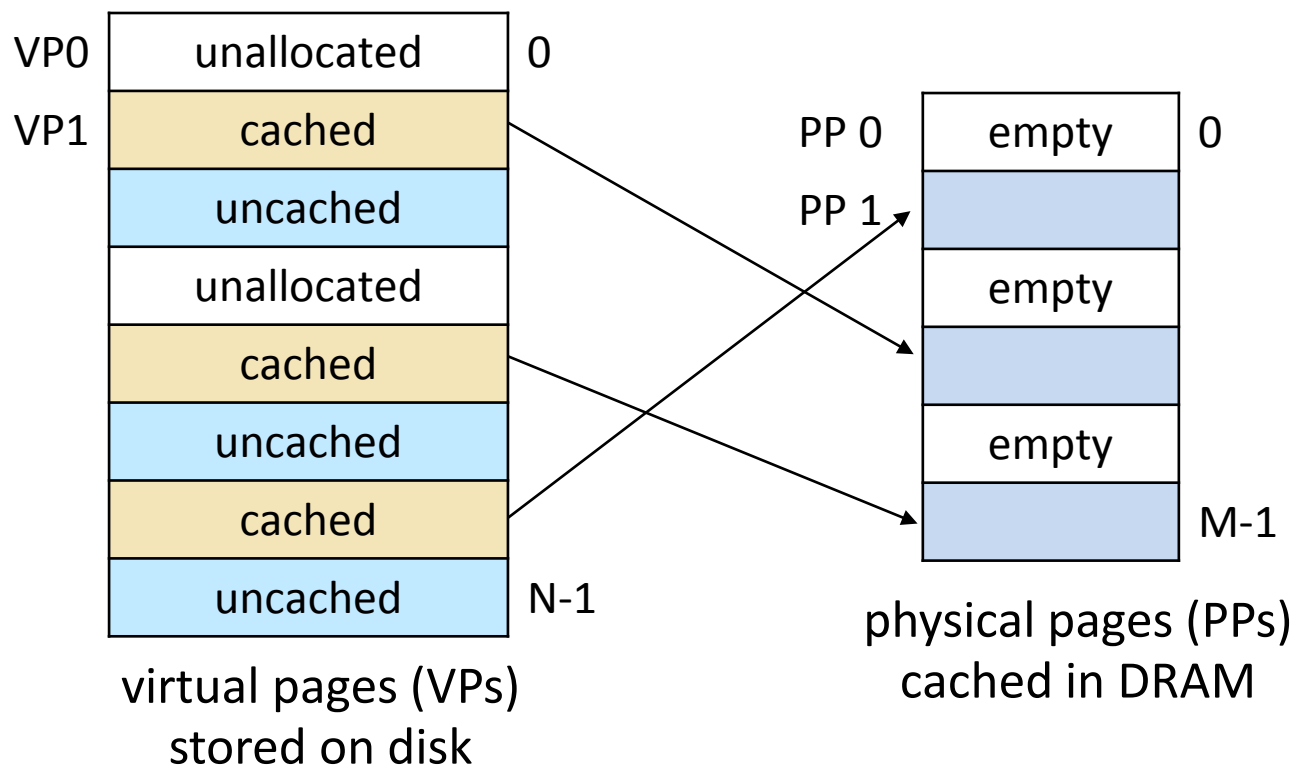
- ***Address space***—an ordered set of integer addresses ( $> 0$ )
  - if addresses are consecutive, the address space is *linear*
- ***Virtual address space***— $N = 2^n$  virtual addresses.
  - $n$  is the number of bits needed to represent the largest address
  - typically, modern virtual address spaces are 32-bit or 64-bit
- ***Physical address space***—corresponds to the  $M = 2^m$  bytes of physical memory in the system.

# Main Memory as a Cache

- Think of virtual memory as an array of  $N$  contiguous byte-sized cells stored on disk.
  - each byte has a unique VA that is an index into the array

- Data on disk is partitioned into blocks (called pages) that serve as the transfer units.

- page size  $P=2^p$
- How many VPs?
- How many PPs?



# DRAM Cache Organization

---

- Misses in DRAM caches are much more expensive than those in SRAM caches.
  - DRAM is  $\sim 10$  times slower than SRAM
    - DRAM:  $\sim 10$  ns / SRAM:  $\sim 1$  ns
  - Spinning disk is  $\sim 1\text{M}$  times slower than DRAM
    - Disk:  $\sim 10$  ms / DRAM:  $\sim 10$  ns
    - Cost of reading the first byte from a (spinning) disk sector is  $\sim 100\text{K}$  times slower than reading successive bytes in the sector
  - Read cost basically uniform on SSD, still gap of  $\sim 10\text{k}$ 
    - SSD:  $\sim 100$  microseconds / DRAM:  $\sim 10$  ns

# DRAM Cache Organization

---

- The organization of DRAM caches is driven by ***the large miss penalty***
  - DRAM caches are fully associative
  - virtual pages are large (4-8 KB)

# Page Table

---

- If a VP is cached in DRAM, which physical page is it cached in?
- If there is a miss:
  - Where is the virtual page stored on disk?
  - Which page in physical memory is the victim?

# Page Table

---

- ***Page table***
  - Data structure stored in physical memory that maps virtual addresses to physical addresses
- Address translation hardware reads the page table
- The OS maintains the contents of the page table and transfers pages between disk and DRAM.

# Page Table Entries

---

- A page table is array of *page table entries* (PTEs).
- Each VP in the virtual address space has a PTE at a fixed offset in the page table.




# Page Table Entries

---

- Assume that each PTE consists of a *valid bit* and a *k*-bit *address field*.
- Valid bit indicates if the VP is currently cached in DRAM
  - if *valid bit* = 1
    - Address field → PP in DRAM
  - if *valid bit* = 0
    - Null address → unallocated VP
    - Non-null address → start of VP on disk

# Example: Page Table




# Page Hit

---

- The ***memory management unit*** (MMU) is dedicated hardware on the CPU that uses the virtual address to locate and read data from memory

# Example: Page Table



*What happens on a read to VP 2?*

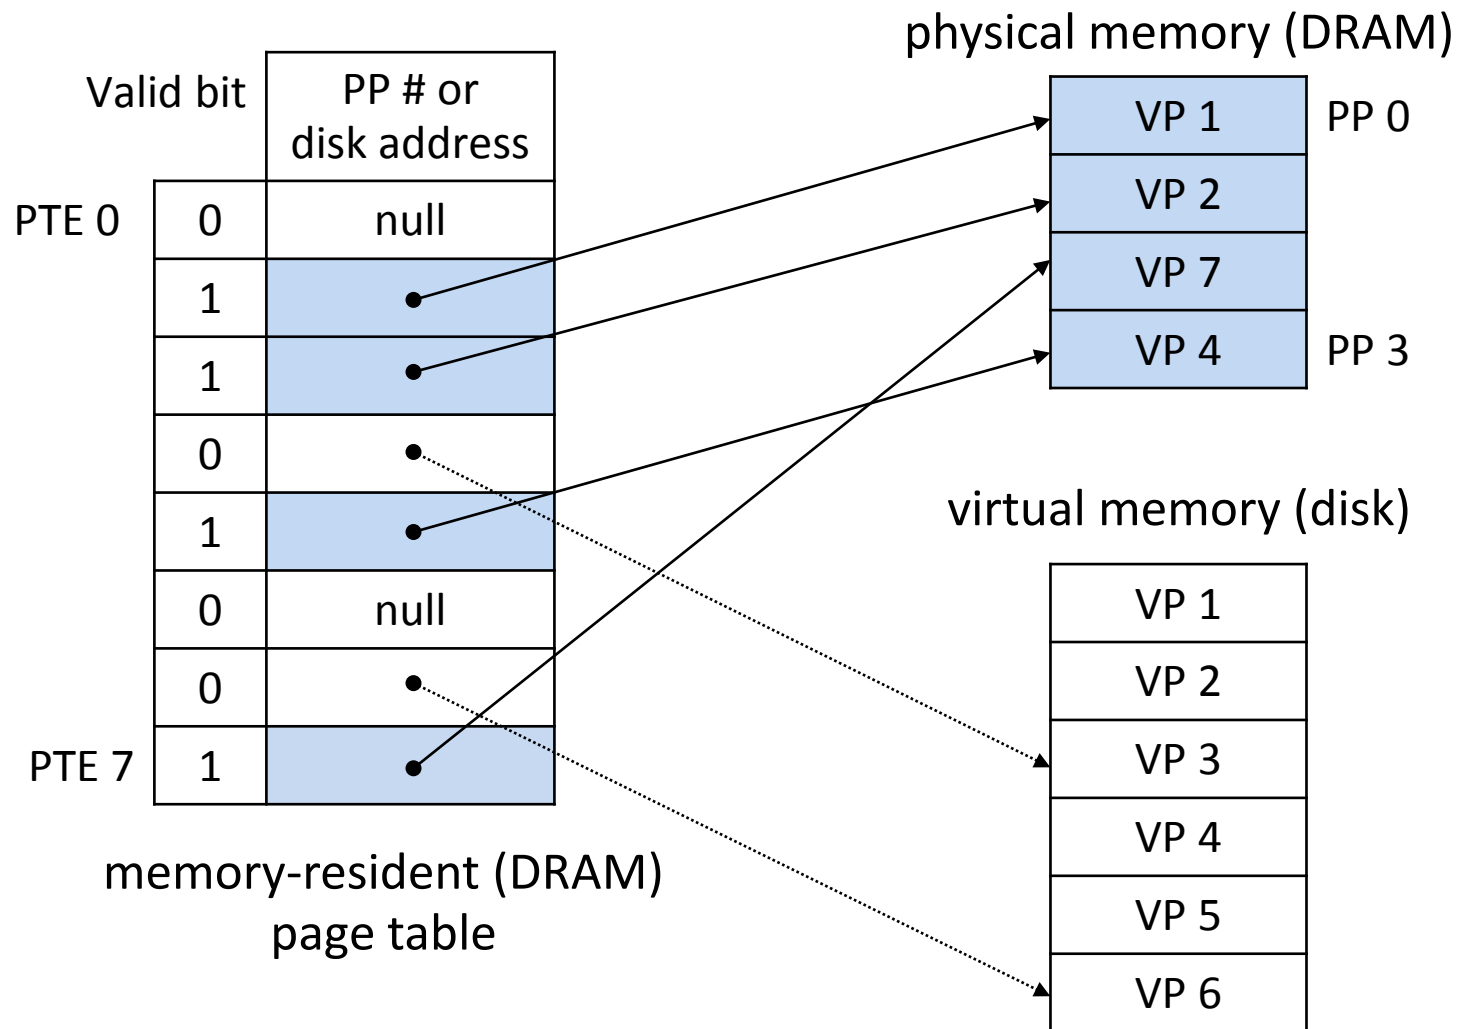
# Page Hit

---

CPU reads a word of virtual memory contained in VP 2 (cached in DRAM):

- Valid bit of PTE 2 is set
  - MMU knows VP 2 is cached, a ***Page Hit*** occurs
- MMU gets address of the cached page (PP 1) from the address field of PTE 2

# Example: Page Table



*What happens on a read to VP 3?*


# Page Fault

---

When the CPU reads a word of VM contained in VP 3:

- The valid bit in PTE 3 is not set
  - Indicates to MMU that VP 3 is uncached, a ***page fault*** occurs
  - Triggers a ***page fault exception*** in kernel
- The page fault exception handler selects a victim page (say, VP 4) and swaps.
- The kernel modifies PTE 4 and PTE 3
- At handler return, faulting instruction restarts, resulting in a page hit

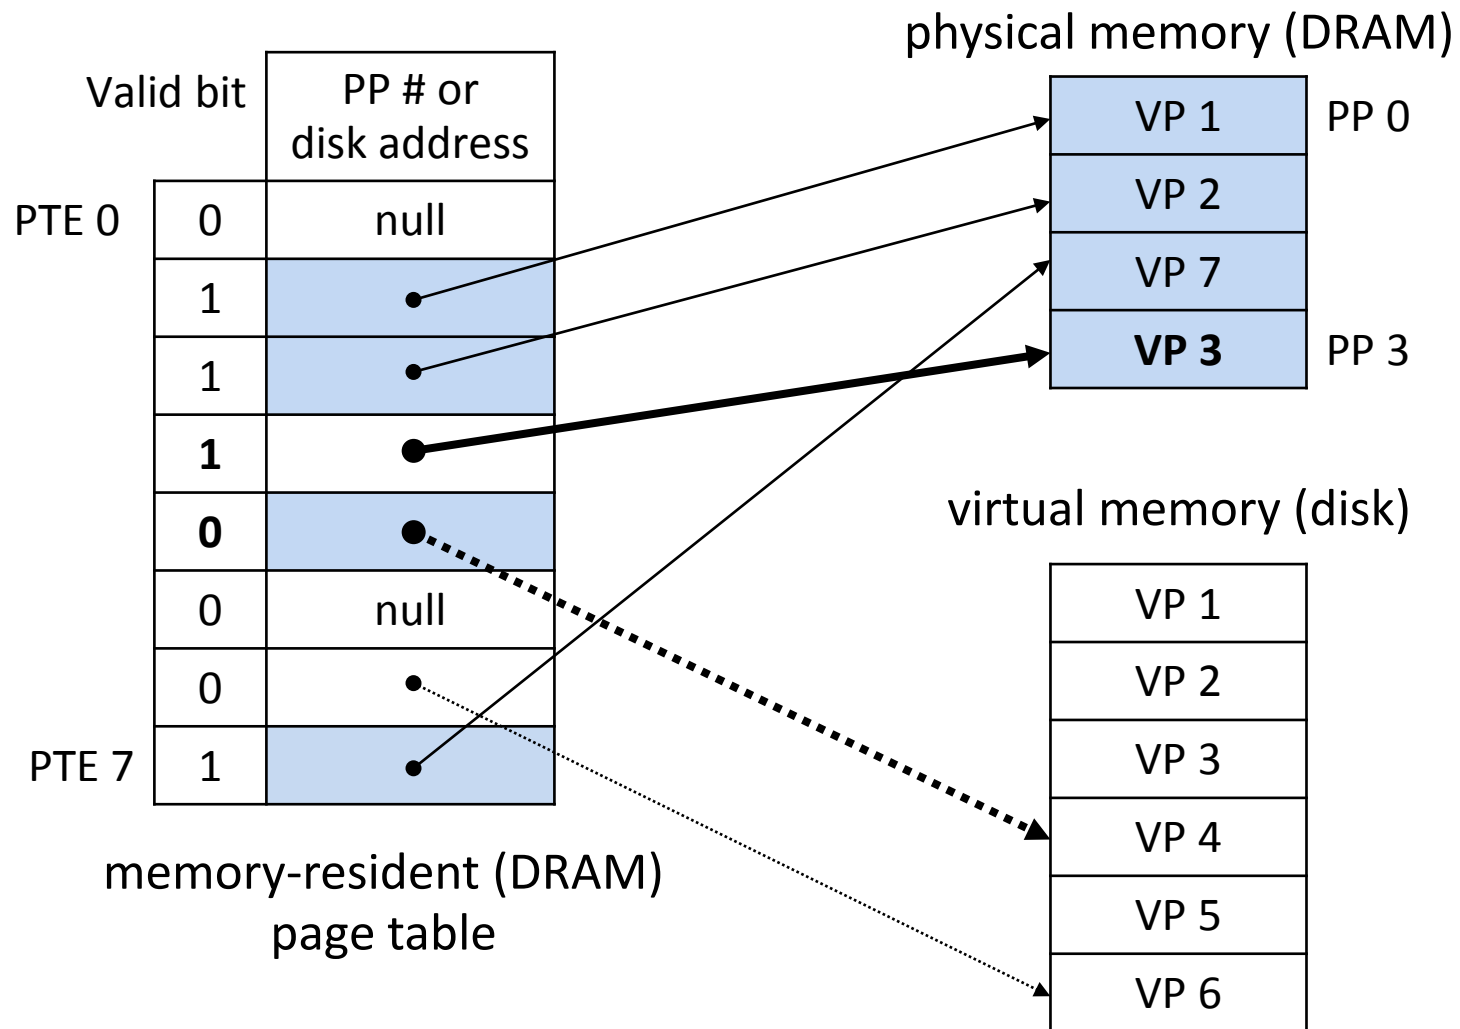
# Example: Page Table



*Read of VP 3 caused page fault*



# Example: Page Table



*VP 3 replaced VP 4 in DRAM (PP 3)*

# Memory Request Actions

---

1. Virtual address sent to the MMU
2. The MMU gets page table entry (PTE) address
  - requests PTE from the cache/main memory
3. The cache/main memory returns the PTE to the MMU.

# Page Hit Actions

---

4. The MMU constructs the physical address
  - Sends it to the cache/main memory
5. Cache/main memory returns the requested data

# Page Fault Actions

---

When (PTE valid=0):

4. The MMU triggers a page fault exception, transferring control to fault handler in OS
5. The fault handler identifies a victim page in physical memory (pages out to disk if needs write-back)
6. The fault handler pages in the new page and updates the PTE in memory
7. The fault handler returns to original process, restarting the faulting instruction—CPU resends VA, page hit

# VM Simplifies . . .

---

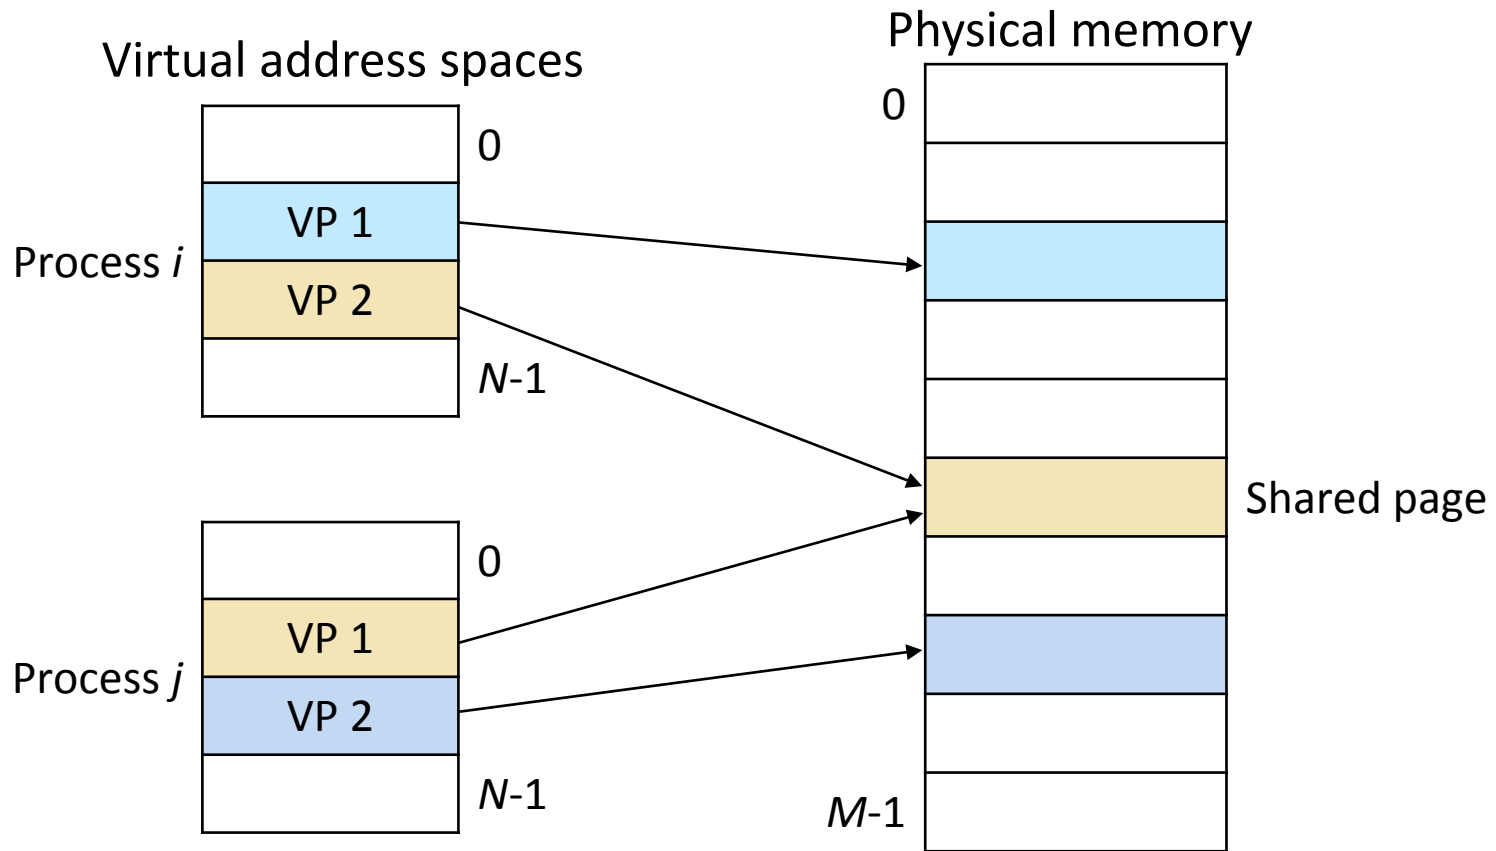
- *Linking*—each process uses the same basic format for its memory image, regardless of where code and data actually reside in physical memory (allows uniformity)
  - `.text` always starts at virtual address `0x08048000`

# VM Simplifies . . .

---

- *Sharing*—provides a consistent mechanism for processes to share code and/or data between processes
  - mapping appropriate VPs in different processes to the same PP

# Memory Management



The OS provides a separate page table, and thus a separate virtual address space, for each process.

# VM Simplifies . . .

---

- *Memory allocation*—simple mechanism for allocating additional memory to user processes (heap space).
  - allocate  $k$  contiguous VPs, no need for PPs to be contiguous



# Memory Protection

---

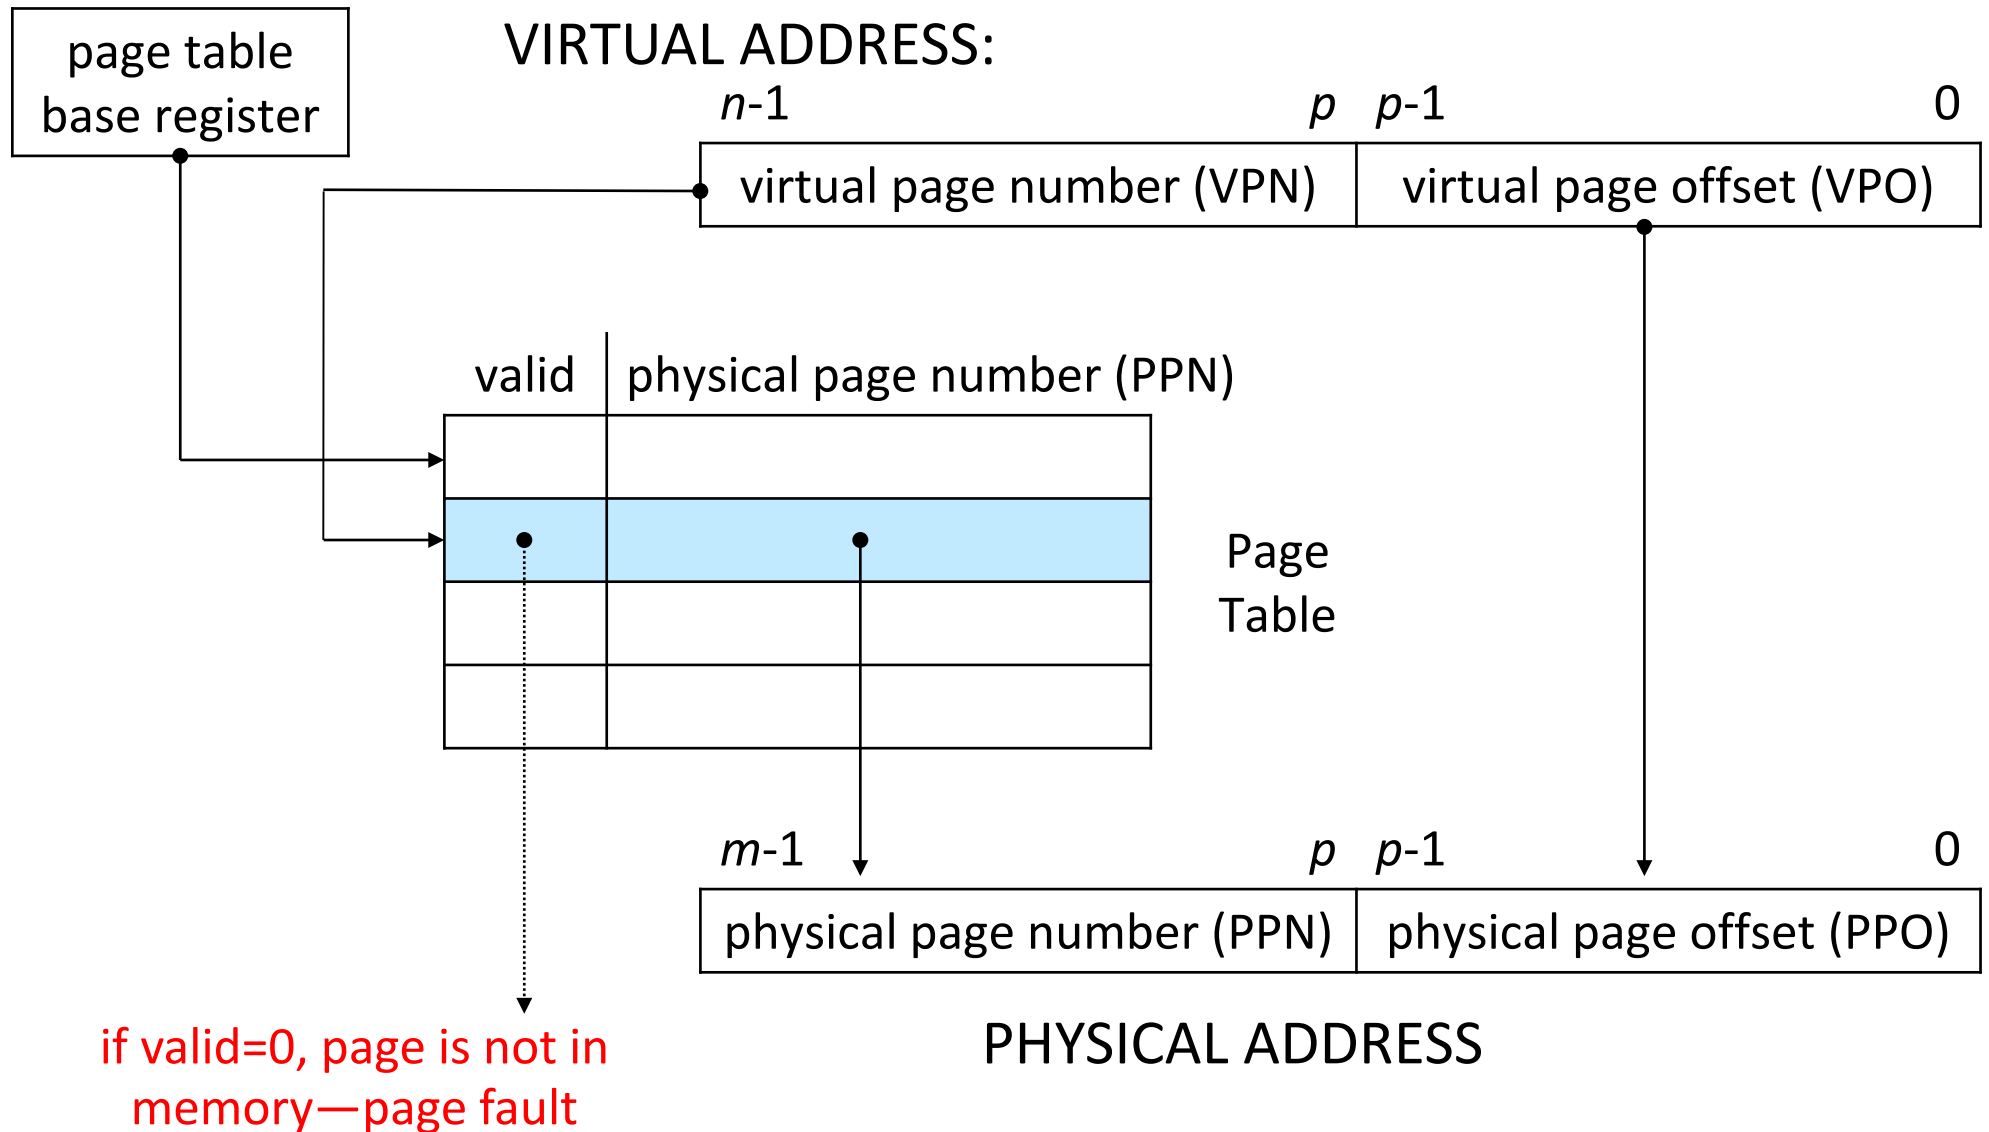
- The OS controls access to the memory system
- A user process should be prevented from:
  - modifying its read-only `.text` section
  - reading or modifying any code/data in the kernel
  - reading or modifying the private memory of other processes
  - modifying any VPs shared with other processes (unless all parties explicitly allow it)

# Memory Protection

---

- The OS controls access to the memory system
- Accomplished by adding permission bits to PTE
  - Indicate a process's read/modify access
  - if an instruction violates these permissions, CPU triggers a general protection fault (typically “segmentation fault”)

# Address Translation



# Exercise: Page Sizes

---

- Suppose 32-bit virtual addresses, 24-bit physical addresses.
- Page size  $P = 1$  KB
  - $\log_2 1024 = 10$ , 10 virtual page offset (VPO) bits
  - $2^{32} / 2^{10} = 2^{22}$ , 22 virtual page number (VPN) bits
  - $\log_2 1024 = 10$ , 10 virtual page offset (PPO) bits
  - $2^{24} / 2^{10} = 2^{14}$ , 14 physical page number (PPN) bits

# Exercise: Page Sizes

---

- Suppose 32-bit virtual addresses, 24-bit physical addresses.
- Page size  $P = 4$  KB
  - VPO bits?
  - VPN bits?
  - PPO bits?
  - PPN bits?

# Translation Lookaside Buffer


---

- For every virtual address generated, the MMU must refer to a PTE to get the corresponding physical address
  - If the PTE is cached in L1, this costs a few cycles
  - If not, it costs tens to hundreds of cycles (L2 or main memory)

# Translation Lookaside Buffer

---

- A ***translation lookaside buffer*** (TLB), small cache of PTEs in the MMU, can eliminate this cost.
- TLB is virtually addressed and each block holds a PTE
  - typically has a high degree of associativity,  $T=2t$  sets



# TLB Hit Actions

---

1. CPU generates virtual address
2. MMU looks up PTE in the TLB
3. TLB hit:
  4. Fetch the PTE from the TLB
5. Construct physical address and send to cache/main memory
6. Cache/main memory returns the requested data



# TLB Miss Actions

---

1. CPU generates virtual address
2. MMU looks up PTE in the TLB
3. TLB miss:
  4. PTE requested from the cache/main memory
  5. TLB is updated with new PTE (evicting another PTE from the TLB)
6. Construct physical address and send to cache/main memory
7. Cache/main memory returns the requested data