

# CS 4400

# Computer Systems

---

## LECTURE 22

Concurrent programming, threads, and  
shared variables

(Chapter 12.3-12.7)

# Application-Level Concurrency

---

- Computing in parallel on multicores
  - logical vs. physical concurrency
- Accessing slow I/O devices
  - already done by kernel, can be done at app-level
- Interacting with humans
  - create a separate concurrent flow to handle each user action
- Reducing latency by deferring work
  - defer work to a concurrent flow that runs at low priority
- Service multiple network clients
  - create a separate concurrent flow for each client (more later)

# Building Concurrent Programs


---

- Processes
  - Each logical control flow is a processes (kernel-scheduled)
  - Separate virtual address spaces—use IPC primitives to communicate
- Threads
  - Logical flows that run in the context of a single process
  - Kernel schedules each thread
  - Hybrid approach—kernel-scheduled, shared address space
- Event-driven
  - Single non-blocking process / thread


# Concurrency w/ Processes

---


- Accept client connection requests in parent, and then create a new child process to service each new client.
- Example: two clients, one server listening



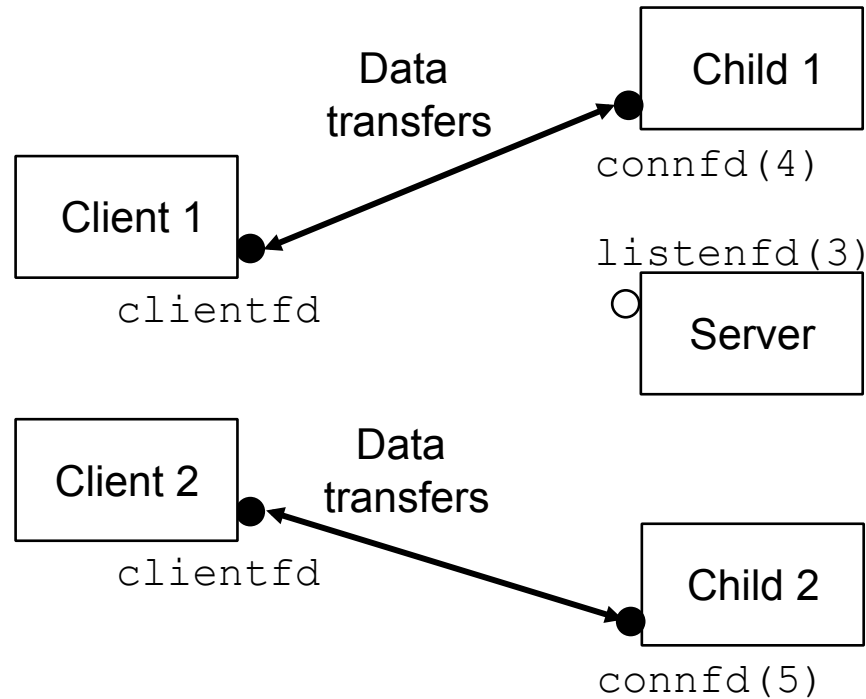
Step 1: Server accepts connection request from client



Step 2: Server forks a child process to service the client.



Step 3: Server accepts another connection request.



*Step 4:* Server forks another child to service new client.

(Parent is waiting for next connection request and two children are servicing their respective clients concurrently.)

# Pros and Cons of Processes

---

- **Pro:** Clean model for sharing state information between parents and children
  - file tables are shared (child gets copy of socket descriptors)
  - user address spaces are not shared (cannot overwrite virtual memory of another process)
- **Con:** Separate address spaces make it more difficult for processes to share state information.
  - must use explicit *interprocess communications (IPC)* mechanisms
- **Con:** Performance as good



# Threads

---

- A ***thread*** is a logical flow that runs in the context of a process.
  - So far, our programs have consisted of a single thread
  - The kernel automatically schedules threads
- Each thread has its own ***thread context***
  - thread ID (TID)—a unique integer
  - stack and stack pointer
  - program counter, gen-purpose registers, and condition codes

# Threads

---

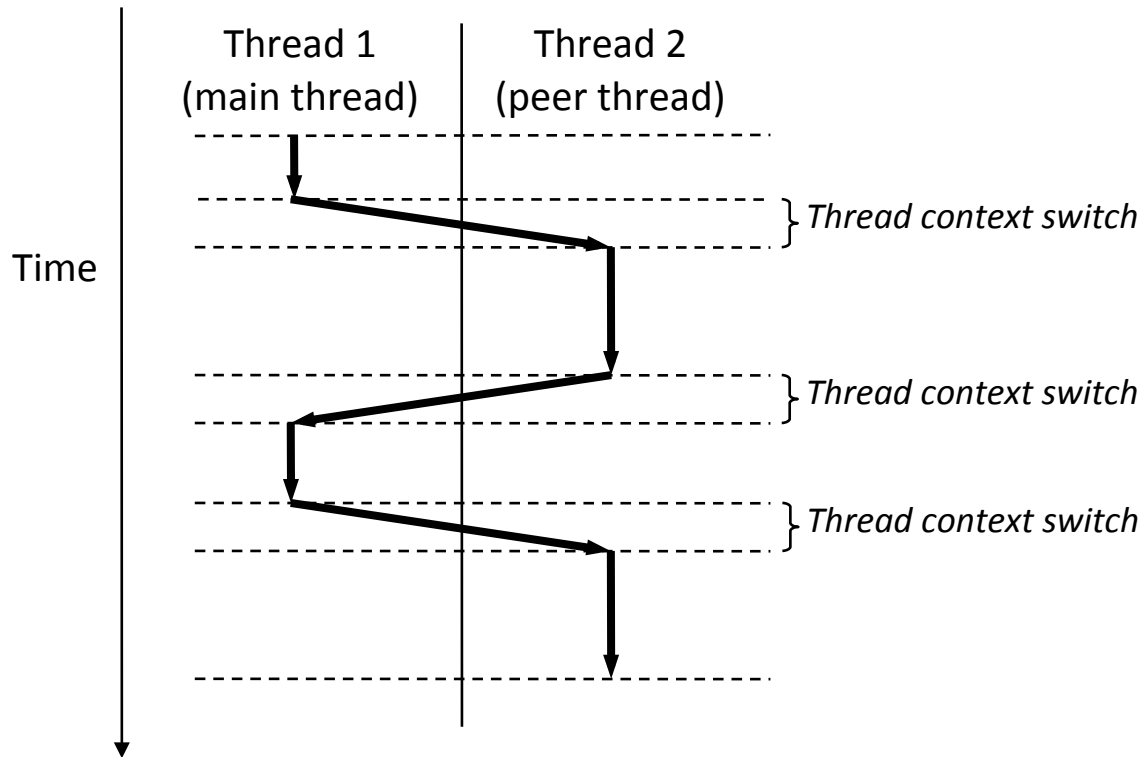
- All threads running in a process share the entire virtual address space sharing
  - code, read/write data, the heap, any shared library code/data, and the set of open files
  - if a shared memory location is modified by one thread, the other threads see the change (if they read the memory loc)

# Threads vs. Processes

---

- A thread context switch can be faster than a process context switch
  - a thread context is much smaller than a process context
- Threads are not organized in a parent-child hierarchy
  - threads associated with a process form a pool of peers, independent of which threads were created by which other threads
  - a thread can kill any of its peers, or wait for any of its peers to terminate
  - each peer can read or write the same shared data

# Execution Model



Control passes to the peer thread because the main thread executes a slow system call or is interrupted by the system's interval timer.

- Each process begins life as a single, main thread
- The main thread creates a peer thread, and from that point the two threads run concurrently

# Creating Threads

---

```
typedef void* (func) (void*);
```

```
int pthread_create(pthread_t* tid,  
    pthread_attr_t* attr, func* f, void* arg);
```

- Creates a new thread and runs the thread routine `f` in the context of the new thread and with input argument `arg`.
- `attr` can be used to change the default thread attributes.
  - we'll always use `NULL`
- Upon return, `tid` is set to the (opaque) ID of the thread
- A thread can determine its own thread ID using:

```
pthread_t pthread_self(void);
```

# Terminating Threads

---

- A thread terminates in one of the following ways.
  - Its top-level thread routine returns
  - Through (if called by the main thread, it waits for all peer threads):

```
int pthread_exit(void* thread_return);
```

- Calls `exit`, which terminates the process and all associated threads.
- Another peer thread calls `pthread_cancel` with the ID of the current thread.

```
int pthread_cancel(pthread_t tid);
```

# Reaping Terminated Threads

---

```
int pthread_join(pthread_t tid,  
                 void** thread_return);
```

- Blocks until thread `tid` terminates.
- Assigns the `void*` returned by the thread routine to the location pointed to by `thread_return`.
- Reaps any memory resources held by the terminated thread.
- Unlike `wait_pid`, this function can only wait for a specific thread to terminate.

# Example: Pthreads

---

```
/* Pthreads is a standard interface for manipulating
 * threads from C programs. */

#include "csapp.h"
void* thread(void* vargp);

/* main thread */
int main() {
    pthread_t tid;    /* thread ID of peer thread */
    /* create peer thread */
    Pthread_create(&tid, NULL, thread, NULL);
    /*--Now, main thread and peer thread are running concurrently.--*/
    /* wait for peer thread to terminate */
    Pthread_join(tid, NULL);
    /* terminate all threads */
    exit(0);
}

/* The code and local data for a thread are encapsulated in
 * a thread routine. Each thread routine takes as input a
 * single generic pointer and returns a generic pointer. */
void* thread(void* vargp) {
    printf("Hello, world!\n");
    return NULL;    /* terminate peer thread */
}
```



# Detaching Threads

---

- At any time, a thread is joinable or detached
  - **joinable**—the thread can be reaped and killed by other threads, at which time its memory resources are freed
  - **detached**—the thread cannot be reaped or killed by other threads, and its memory resources are free automatically by the system when it terminates
- By default, all threads are created joinable
- To avoid memory leaks, each joinable thread should either be reaped by another thread or detached

```
int pthread_detach(pthread_t tid);
```

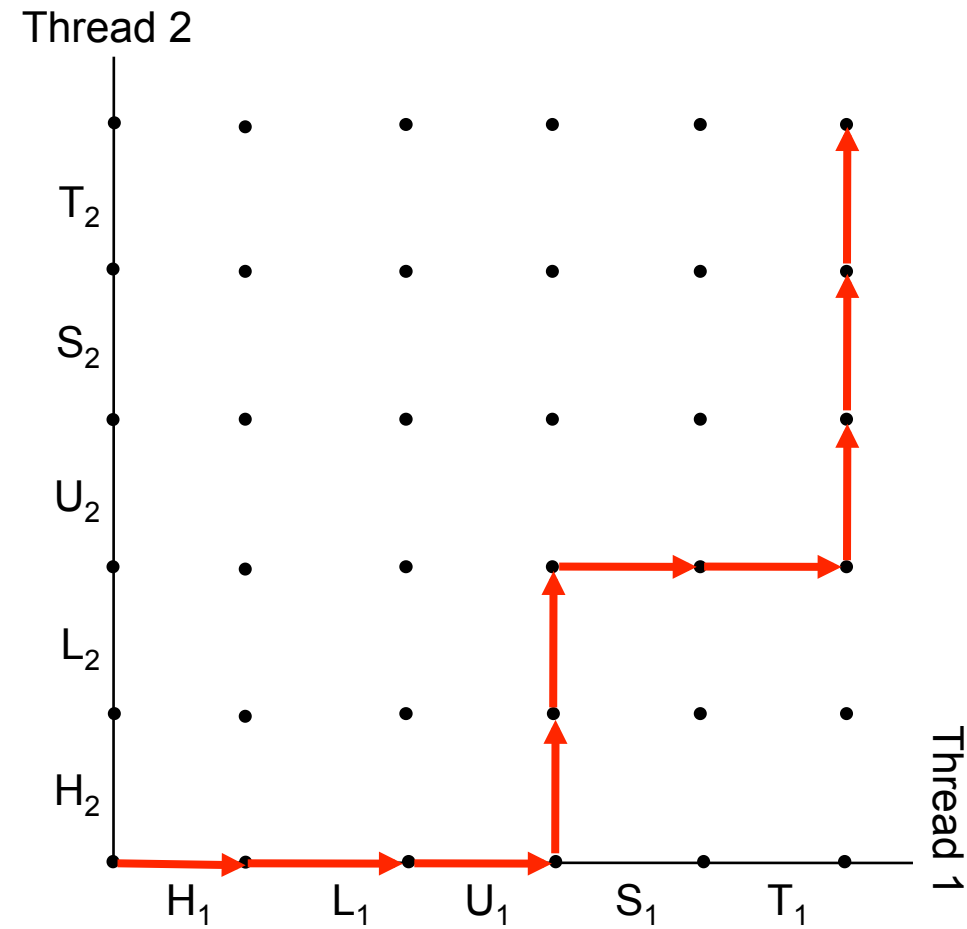
# Mapping Variables to Memory

---

- ***Global variable***—any variable declared outside of a function.
  - one instance that can be referenced by any thread
- ***Local automatic variables***—any variable declared inside a function without the `static` attribute.
  - each thread's stack contains its own instance (even if multiple threads execute the same thread routine)
- ***Local static variables***—any variable declared inside a function with the `static` attribute.
  - like global variables, one instance for all threads

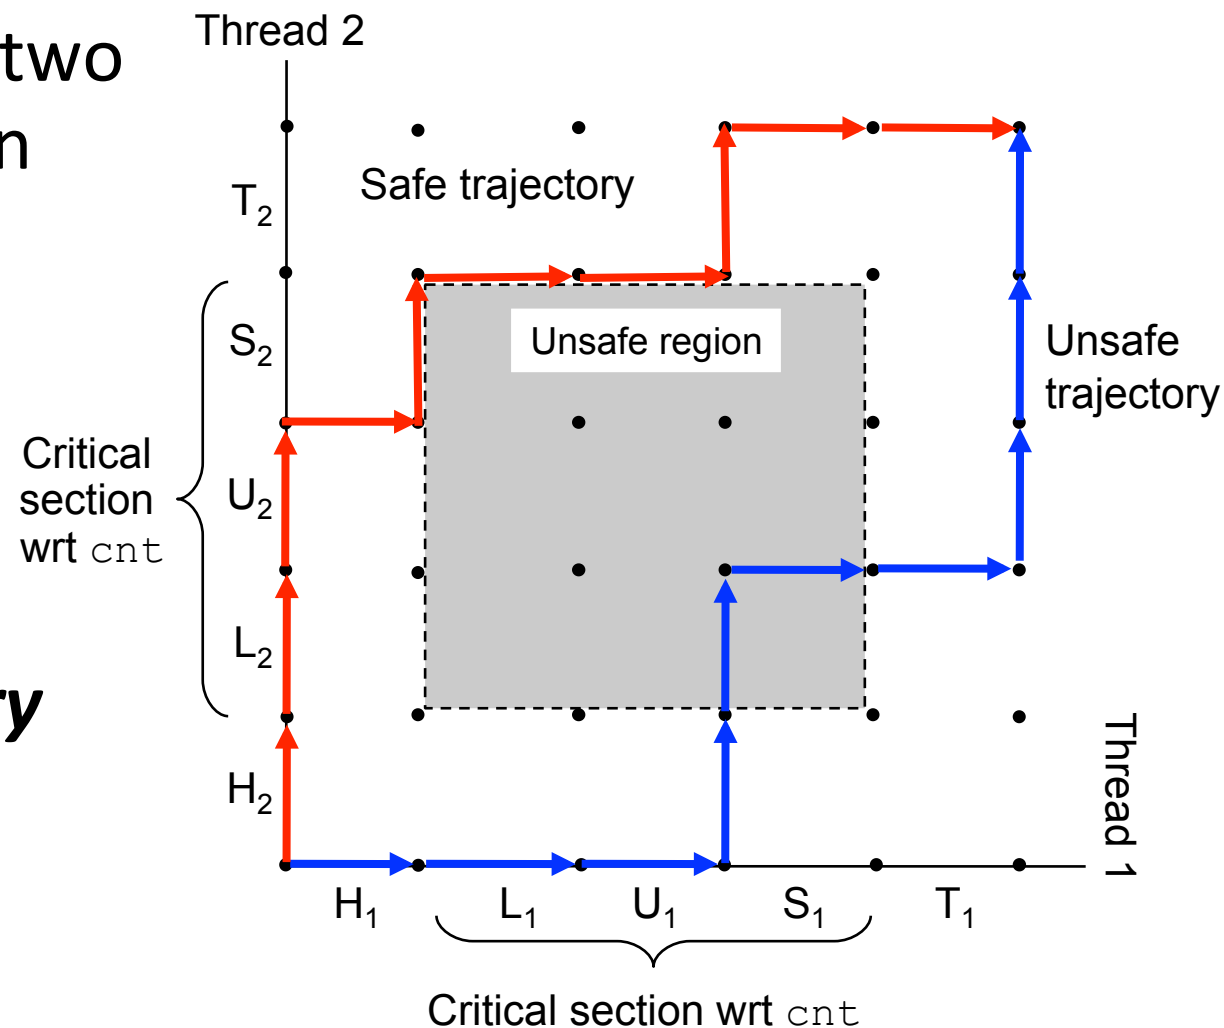
# Process Graph

- Models the execution of  $n$  threads as a trajectory through an  $n$ -dimensional Cartesian space.
- each axis  $k$  shows the progress of thread  $k$
- each point  $(l_1, l_2, \dots, l_n)$  represents the state where thread  $k$  has completed instruction  $l_k$
- the trajectory corresponds to the ordering of instructions



# Critical Section

- Instructions  $L_i$ ,  $U_i$ , and  $S_i$  constitute a ***critical section*** for thread  $i$ .
- The intersection of two critical sections is an ***unsafe region***.
- A ***safe trajectory*** skirts the unsafe region.
- An ***unsafe trajectory*** touches any part of the unsafe region.



# Semaphore

---

- A global variable  $s \geq 0$  that can only be manipulated using one of two operations: P and V.
- P(s)
  - if  $s \neq 0$ ,  $s--$  and return (occurs indivisibly)
  - if  $s = 0$ , suspend the process until  $s$  becomes nonzero (process is restarted by a V operation), after restarting  $s--$  and return
- V(s)
  - $s++$  and check to see if any processes are blocked in a P operation waiting for  $s$  to become nonzero (restarts exactly one of such processes)
  - increment occurs indivisibly

# Posix Semaphores

---

- Functions for manipulating semaphores.

```
int sem_init(sem_t* sem, 0, unsigned int value);  
int sem_wait(sem_t* sem);           /* P(s) */  
int sem_post(sem_t* sem);          /* V(s) */
```

- Example:

```
sem_t mutex; /* semaphore to synch cnt access */  
sem_init(&mutex, 0, 1);           /* init mutex */  
...  
for(i = 0; i < NITERS; i++) {  
    sem_wait(&mutex);             /* protect shared */  
    cnt++;                        /* variable cnt */  
    sem_post(&mutex);  
}
```

# Producer-Consumer Model

---

- Producer and consumer threads share a bounded buffer, with  $n$  slots.
  - producer thread adds items to the buffer
  - consumer thread retrieves items from the buffer
- Must guarantee mutually-exclusive access to the buffer, and that the producer/consumer cannot access the buffer if it is full/empty.

```
typedef struct {
    int* buf;      /* Buffer array */
    int n;         /* Max # of slots */
    int front;     /* buf[(front+1)%n] is 1st item */
    int rear;      /* buf[rear%n] is last item */
    sem_t mutex;   /* Protects accesses to buf */
    sem_t slots;   /* Counts available slots */
    sem_t items;   /* Counts available items */
} sbuf_t;
```

# Other Concurrency Issues

---

- We've looked at techniques for mutual exclusion and producer-consumer synchronization, a small part of concurrent programming.
- Synchronization is a fundamentally difficult problem that raises issues that do not arise in sequential programs.
- What follows is a sample of the issues programmers must be aware of when writing concurrent programs.
- Presented in the context of threads, the issues exist whenever concurrent flows manipulate shared resources.



# Thread Safety

---

- A function is ***thread-safe*** iff it always produces correct results when called repeatedly from multiple concurrent threads.
  - a function that is not thread-safe is called thread-unsafe
- Four (non-disjoint) classes of thread-unsafe functions:
  - Class 1: functions that do not protect shared variables
  - Class 2: functions that keep state across multiple invocations
  - Class 3: functions that return a pointer to a static variable
  - Class 4: functions that call thread-unsafe functions

# Class 1: Shared Variables

---

```
/* thread-unsafe routine */  
void* count(void* arg) {  
    int i;  
    for(i = 0; i < NITERS; i++)  
        cnt++;  
    return NULL;  
}
```

- To make thread-safe, protect the shared variable with synchronization operations.
- **Pro:** No changes in the calling program required.
- **Con:** Synchronization operations will slow down the function even when called from single-threaded program

# Class 2: Keeps State Across Calls

---

```
unsigned int next = 1;
/* rand - return pseudo-random integer on 0..32767 */
int rand(void) {
    next = next*1103515245 + 12345;
    return (unsigned int) (next/65536) % 32768;
}
/* srand - set seed for rand() */
void srand(unsigned int seed) {
    next = seed;
}
```

- Calling `rand` repeatedly from a single thread is correct.
  - What can happen if it is called from multiple threads?
- To make thread-safe, we must rely on the caller to pass state information via arguments.
  - forces a change in the code of the calling routine
  - potentially 100s of call sites, a difficult and error-prone change

# Class 3: Returns Pointer to Static

---

- Some functions compute a result in a local `static` variable and return a pointer to that variable.
  - results being used by one thread may be silently overwritten by another thread
- To make thread-safe, require the caller to pass the address of the variable in which to store the result.
  - removes shared variable, requires change in calling code
- Another option is the ***lock-and-copy*** technique.
  - associate a mutex with the thread-unsafe function
  - especially useful when the thread-unsafe function is impossible to modify (e.g., it is linked from a library)

# Lock-and-Copy

---

- At each call site:
  - dynamically allocate memory for the result
  - lock the mutex
  - call the thread-unsafe function
  - copy the result returned by the function to this memory
  - unlock the mutex

```
struct hostent* gethostbyname_ts(char* hostname) {
    struct hostent *sharedp, *unsharedp;

    unsharedp = Malloc(sizeof(struct hostent)); /* dyn mem */
    P(&mutex);                                /* lock mutex */
    sharedp = gethostbyname(hostname); /* thread-unsafe fn */
    *unsharedp = *sharedp;                /* copy to private struct */
    V(&mutex);                                /* unlock mutex */
    return unsharedp;
}
```

# Class 4: Calls Thread-Unsafe

---

- If function  $f$  calls thread-unsafe function  $g$ ,  $f$  may or may not also be thread-unsafe.
- If  $g$  keeps state across multiple invocations, then  $f$  is also thread-unsafe.
  - only solution is to rewrite  $g$
- If  $g$  does not protect shared variables or returns a pointer to a static variable,  $f$  may still be thread-safe.
  - solution is to protect call to  $g$  with a mutex (like previous example)

# Reentrancy

---

- Reentrant functions do not reference any shared data when they are called by multiple threads.
- The set of reentrant functions is a proper subset of the thread-safe functions.
  - due to the lack of synchronization ops, reentrant functions are typically more efficient than non-reentrant thread-safe functions
- The only way to convert a Class 2 thread-unsafe function into a thread-safe one is to rewrite it to be reentrant.

```
/* rand_r - a reentrant pseudo-random integer generator */  
int rand_r(unsigned int* nextp) {  
    *nextp = *nextp * 1103515245 + 12345;  
    return (unsigned int)(*nextp / 65536) % 32768;  
}
```

# Determining Reentrancy

---

- ***Explicitly reentrant***—all function arguments are passed by value and all data references are to local automatic stack variables.
- ***Implicitly reentrant***—allows some parameters in an otherwise explicitly-reentrant function to be pointers.
  - thus, it is a reentrant function only if the calling threads are careful to pass pointers to non-shared data
  - example: function `rand_r`
- Why is function `gethostbyname_ts` thread-safe, but not reentrant?



# Races

---

- A ***race*** occurs when the correctness of a program depends on one thread reaching point ***x*** in its control flow before another thread reaches point ***y***.
- Threaded programs must work correctly for any feasible trajectory.
  - Often programmers assume that threads will take a particular trajectory through the execution state space.

# Deadlock

---

- A run-time error where a collection of threads are blocked, waiting for a condition that will never be true.
- The programmer has incorrectly ordered the semaphore ops.
- In state d, each program is waiting for the other to do a V op that won't occur.

# Avoiding Deadlock

---

- Deadlock is difficult to predict in a program.
  - some trajectories will skirt the deadlock region
  - others will be trapped by it
- When semaphores are used for mutual exclusion, a simple rule can be applied.
  - A program is deadlock-free if, for each pair of mutexes ( $s$ ,  $t$ ) in the program, each thread that holds both  $s$  and  $t$  simultaneously locks them in the same order.
- In our example, lock  $s$  first then  $t$ , in each thread.

# Summary

---

- A concurrent program consists of a collection of logical flows that overlap in time.
  - via processes—scheduled by the kernel, separate address space
  - via threads—scheduled by the kernel, shared address space
- P and V operations on semaphores help to synchronize concurrent accesses to shared data.
  - provides mutually exclusive access to shared data
  - schedules access to shared buffers in producer-consumer programs
- Difficult concurrency issues:
  - thread safety, reentrant functions, races, deadlocks